

Exercise3_Group1

February 19, 2025

1 Exercise 3: Data-Driven Computing Architectures

In this exercise, we will work with Delta tables and the Medallion Architecture. You can gain 15 points, which will be awarded based on how effectively you implement the Bronze, Silver, and Gold layers, the quality of your data transformations and analysis, and how well your visualizations communicate insights. A well-designed pipeline should also allow new files to be uploaded and processed smoothly without requiring major modifications.

Make sure to include clear explanations of what you did and why throughout your report and visualizations. Grading will also consider how well you justify your choices, so do not just present results but explain your reasoning.

Useful links: - <https://docs.delta.io/latest/index.html> - <https://delta.io/blog/delta-lake-medallion-architecture/>

1.1 Scenario

You're a Data Engineer at a manufacturing company that produces industrial components. The factory runs 10 specialized machines, producing drill bits, gears, shafts, conveyor belts, turbine parts, robot components, stamped metal, polished surfaces, laser-cut materials, and 3D-printed prototypes.

Each machine is equipped with sensors to track performance, while production logs record output and defects, and maintenance records document repairs. A team of production operators manages manufacturing, while maintenance operators handle scheduled and emergency repairs.

Recently, management has raised concerns about machine efficiency, defect rates, and maintenance costs, and they want continuous data-driven insights to improve operations. You have been given raw data from three sources:

- **Sensor Data:** Real-time readings from industrial machines.
- **Production Logs:** Daily records of production output and defects.
- **Maintenance Records:** Logs of scheduled and emergency maintenance events.

1.2 Your Assignment

Your task is to build a Medallion Architecture pipeline using Delta Lake to clean, structure, and analyze this data. 1. **Ingest the raw data** into **Bronze Layer** Delta tables.

2. **Clean and standardize the data** in the **Silver Layer**.

3. **Aggregate and generate business insights** in the **Gold Layer**.
 4. **Visualize key metrics** to make informed decisions, using for example Matplotlib or Seaborn.
-

1.3 Data Description

The csv files are available in the `/shared` folder in Noppe. **## 1. Sensor Data**

Captures real-time sensor readings from machines, tracking temperature, vibration, power consumption, and operational status.

1.3.1 Key Fields:

- `sensor_id`: Unique identifier for the sensor.
 - `machine`: Name of the machine.
 - `time_stamp`: Timestamp of the sensor reading.
 - `temp C`: Temperature reading (°C).
 - `vibration_lvl`: Vibration level reading.
 - `power_kW`: Power consumption (kW).
 - `def_ct`: Number of defective sensor readings.
 - `status_flag`: Operational status (e.g., “Running”, “Stopped”).
 - `noise_val`: Random noise factor in the data.
 - `extra_param`: Additional machine-related parameter.
-

1.4 2. Production Logs

Tracks machine output, defect rates, and operator activity. Machines in poor condition tend to produce more defects.

1.4.1 Key Fields:

- `log_id`: Unique identifier for the production log.
- `product_type`: Type of product produced.
- `units_produced`: Number of units produced.
- `defective_units`: Number of defective units.

- **time_stamp:** Timestamp of production record.
 - **machine:** Machine responsible for production.
 - **operator_id:** Identifier of the operator overseeing production.
 - **remarks:** Additional notes about the production process (e.g., quality concerns, machine adjustments).
 - **batch_info:** Batch identifier for tracking specific production runs.
-

1.5 3. Maintenance Records

Logs maintenance activities, including scheduled upkeep, emergency repairs, and associated costs. Machines in poor condition require more frequent emergency maintenance.

1.5.1 Key Fields:

- **maintenance_id:** Unique identifier for the maintenance event.
 - **machine:** Machine undergoing maintenance.
 - **maintenance_date:** Timestamp of maintenance event.
 - **maintenance_type:** Type of maintenance (Scheduled, Unscheduled, Emergency).
 - **duration_minutes:** Length of the maintenance event.
 - **cost:** Cost of the maintenance.
 - **operator:** Identifier of the maintenance operator performing the task.
 - **notes:** Description of the maintenance issue or action taken.
-

1.6 4. Operator Dimension Table (Predefined Silver Table)

A reference table with details about production and maintenance operators.

1.6.1 Key Fields:

- **operator_id:** Unique identifier for the operator.
 - **operator_name:** Full name of the operator.
 - **operator_type:** “Production” or “Maintenance”.
-

1.7 Medallion Architecture Implementation

1.7.1 1. Bronze Layer – Raw Data Storage (2p)

- Ingest raw data as-is into Delta tables.
- No transformations at this stage.

1.7.2 2. Silver Layer – Cleaning & Standardization (3p)

For example: - Convert columns into a proper format.

- Rename columns for consistency, for example (`time_stamp` → `timestamp`).
- Remove duplicate records.

1.7.3 3. Gold Layer – Business Insights (4p)

For example: - Aggregate sensor, production, and maintenance data to create daily machine performance metrics.

- Join tables to uncover correlations between sensor readings, defects, and maintenance events.
- Visualize the tables and metrics. —

1.8 Example Visualizations:

1.8.1 Daily Sensor Metrics

- Average temperature, vibration, and power consumption per machine.
- Number of downtime events (`status_flag = "Stopped"`).

1.8.2 Daily Production Metrics

- Total units produced and defective units per machine.
- Production yield: $(\text{total_units_produced} - \text{defective_units}) / \text{total_units_produced}$

1.8.3 Daily Maintenance Metrics

- Number of maintenance events per machine.
- Total maintenance costs per machine.

1.8.4 Advanced Insights

- Correlation analysis between high vibration levels and production defects.
- Identify the most frequent operator per machine per day.
- Estimate energy consumption trends over time.

You are encouraged to explore and define your own insights.

1.9 Example Directory Structure

```
data_lake/
  bronze/
    sensor_data_bronze/      # Raw Sensor Data (Delta table)
    production_data_bronze/  # Raw Production Data (Delta table)
    maintenance_data_bronze/ # Raw Maintenance Data (Delta table)
  silver/
    sensor_data_silver/      # Cleaned Sensor Data (Delta table)
    production_data_silver/  # Cleaned Production Data (Delta table)
    maintenance_data_silver/ # Cleaned Maintenance Data (Delta table)
    dim_operator_silver/     # Operator Dimension Table (Delta table)
  gold/
    gold_machine_performance/ # Aggregated Machine Performance (Delta table)
```

```
[1]: pip install delta-spark==3.0.0
```

```
Collecting delta-spark==3.0.0
  Downloading delta_spark-3.0.0-py3-none-any.whl.metadata (2.0 kB)
Requirement already satisfied: pyspark<3.6.0,>=3.5.0 in /usr/local/spark/python
(from delta-spark==3.0.0) (3.5.1)
Requirement already satisfied: importlib-metadata>=1.0.0 in
/opt/conda/lib/python3.11/site-packages (from delta-spark==3.0.0) (7.1.0)
Requirement already satisfied: zipp>=0.5 in /opt/conda/lib/python3.11/site-
packages (from importlib-metadata>=1.0.0->delta-spark==3.0.0) (3.17.0)
Collecting py4j==0.10.9.7 (from pyspark<3.6.0,>=3.5.0->delta-spark==3.0.0)
  Downloading py4j-0.10.9.7-py2.py3-none-any.whl.metadata (1.5 kB)
  Downloading delta_spark-3.0.0-py3-none-any.whl (21 kB)
  Downloading py4j-0.10.9.7-py2.py3-none-any.whl (200 kB)
    200.5/200.5 kB
5.5 MB/s eta 0:00:00a 0:00:01
Installing collected packages: py4j, delta-spark
Successfully installed delta-spark-3.0.0 py4j-0.10.9.7
Note: you may need to restart the kernel to use updated packages.
```

```
[2]: from pyspark.sql import SparkSession
from delta import configure_spark_with_delta_pip

# Configure the Spark session with Delta support
builder = SparkSession.builder \
    .appName("Exercise1") \
    .config("spark.sql.extensions", "io.delta.sql.DeltaSparkSessionExtension") \
    .config("spark.sql.catalog.spark_catalog", "org.apache.spark.sql.delta.
↪catalog.DeltaCatalog") \
    .config("spark.jars.packages", "io.delta:delta-core_2.12:3.0.0")

# Create the Spark session
```

```
spark = configure_spark_with_delta_pip(builder).getOrCreate()

print("Spark session with Delta Lake configured successfully!")
spark
```

Spark session with Delta Lake configured successfully!

[2]: <pyspark.sql.session.SparkSession at 0x7fa939709550>

1.10 Imports

```
[3]: from pyspark.sql.functions import col, sum, to_date, avg, count, when, isnan
from pyspark.sql.types import DoubleType, IntegerType, TimestampType
import matplotlib.pyplot as plt
import seaborn as sns
import os
```

1.11 1. Bronze Layer – Raw Data Storage

1.11.1 Sensor_data

```
[4]: sensor_data_path = "sensor_data.csv"

# Read the raw sensor data from CSV
sensor_df = spark.read.format("csv") \
    .option("header", "true") \
    .option("inferSchema", "true") \
    .load(sensor_data_path)

# Whitespaces in the columns names were giving error -> all spaces were changed
↳with underlines
sensor_df = sensor_df.select([col(c).alias(c.replace(" ", "_")) for c in
↳sensor_df.columns])

sensor_df.show(5)
sensor_df.printSchema()
```

```
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|      sensor_id|      machine|      time_stamp|temp_C|vibration_lvl|power_kW|def_ct|status_flag|noise_val|extra_param|
+-----+-----+-----+-----+-----+-----+
|e0310580-9228-434...|DrillPress-100|2024-10-01 00:00:00| 48.1|          1.03|20.9|    0|    Stopped|    0.505|          103|
|496a0d42-dd46-441...|  CNC-Mill-200|2024-10-01 00:02:00| 48.0|          0.9|19.25|    0|    Stopped|    0.716|          189|
```



```
production_df = production_df.select([col(c).alias(c.replace(" ", "_")) for c
↳in production_df.columns])
```

```
production_df.show(5)
production_df.printSchema()
```

```
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|          log_id|   product_type|units_produced|defective_units|
time_stamp|  MachineName|      remarks|operator_id|      batch_info|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|831bdeda-388c-4f3...|   Turbine Hub|          96|          5|2024-10-01
10:05:00|   Turbine-500|Minor delays due ...|   OP17|Batch-Turbine-500...|
|41274f57-82a0-4a4...|Polished Surface|          87|          4|2024-10-01
11:59:00|   Grinder-800|Normal operations...|   OP14|Batch-Grinder-800...|
|7a943af7-dc35-4e7...|   Drill Bit|        137|          8|2024-10-01
21:49:00|DrillPress-100|Slight quality co...|   OP20|Batch-DrillPress-...|
|5365b96b-62e9-470...|   Drill Bit|        135|          7|2024-10-01
05:27:00|DrillPress-100|Normal operations...|   OP10|Batch-DrillPress-...|
|e49c5179-038f-40f...|   Motor Shaft|        122|          6|2024-10-01
04:18:00|   CNC-Mill-200|Slight quality co...|   OP10|Batch-CNC-Mill-20...|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
root
|-- log_id: string (nullable = true)
|-- product_type: string (nullable = true)
|-- units_produced: string (nullable = true)
|-- defective_units: string (nullable = true)
|-- time_stamp: string (nullable = true)
|-- MachineName: string (nullable = true)
|-- remarks: string (nullable = true)
|-- operator_id: string (nullable = true)
|-- batch_info: string (nullable = true)
```

```
[7]: # Write raw data to Bronze Layer as a Delta table
bronze_production_path = "data_lake/bronze/production_data_bronze/"
os.makedirs(bronze_production_path, exist_ok=True)

production_df.write.format("delta") \
    .mode("append") \
    .save(bronze_production_path)

print("Raw production data successfully ingested into the Bronze Layer.")
```



```
bronze_production_df = spark.read.format("delta").load(bronze_production_path)
```

Raw production data successfully ingested into the Bronze Layer.

1.11.3 Maintenance Data

```
[8]: maintenance_data_path = "maintenance_data.csv"

# Read the raw production data from CSV
maintenance_df = spark.read.format("csv") \
    .option("header", "true") \
    .option("inferSchema", "true") \
    .load(maintenance_data_path)

# Whitespaces in the columns names were giving error -> all spaces were changed
↳ with underlines
maintenance_df = maintenance_df.select([col(c).alias(c.replace(" ", "_")) for c
↳ in maintenance_df.columns])

maintenance_df.show(5)
maintenance_df.printSchema()
```

```
+-----+-----+-----+-----+-----+
-----+-----+-----+-----+
|      maintenance_id|      machine|
maintenance_date|maintenance_type|duration_minutes|    cost|operator|
notes|
+-----+-----+-----+-----+-----+
-----+-----+-----+-----+
|e0c400c6-e7a8-4b4...|  Conveyor-400|2024-10-02 05:48:00|      Emergency|
140|1312.39|    MT104|Critical failure ...|
|0f2cbacd-52d6-49a...|DrillPress-100|2024-10-03 03:46:00|      Unscheduled|
49| 263.19|    MT105|Replaced small co...|
|48ab287e-4e96-462...|  Conveyor-400|2024-10-03 18:15:00|      Scheduled|
66| 284.63|    MT101|Calibrated machin...|
|b170c9d2-c3fc-4c9...|  RobotArm-600|2024-10-03 03:04:00|      Scheduled|
91| 265.16|    MT103|Replaced worn-out...|
|4f6c00a5-5eaa-425...|  Conveyor-400|2024-10-04 13:04:00|      Scheduled|
99| 212.26|    MT102|Routine check-up ...|
+-----+-----+-----+-----+-----+
-----+-----+-----+-----+
only showing top 5 rows
```

```
root
|-- maintenance_id: string (nullable = true)
|-- machine: string (nullable = true)
|-- maintenance_date: timestamp (nullable = true)
```

```

|-- maintenance_type: string (nullable = true)
|-- duration_minutes: integer (nullable = true)
|-- cost: double (nullable = true)
|-- operator: string (nullable = true)
|-- notes: string (nullable = true)

```

```

[9]: # Write raw data to Bronze Layer as a Delta table
bronze_maintenance_path = "data_lake/bronze/maintenance_data_bronze/"
os.makedirs(bronze_maintenance_path, exist_ok=True)

maintenance_df.write.format("delta") \
    .mode("append") \
    .save(bronze_maintenance_path)

print("Raw maintenance data successfully ingested into the Bronze Layer.")
bronze_maintenance_df = spark.read.format("delta").load(bronze_maintenance_path)

```

Raw maintenance data successfully ingested into the Bronze Layer.

1.12 Silver Layer – Cleaning & Standardization

1.12.1 Sensor Data

Renaming Columns

```

[10]: bronze_sensor_df.columns

```

```

[10]: ['sensor_id',
      'machine',
      'time_stamp',
      'temp_C',
      'vibration_lvl',
      'power_kW',
      'def_ct',
      'status_flag',
      'noise_val',
      'extra_param']

```

```

[11]: silver_sensor_df = (
    bronze_sensor_df.withColumnRenamed("time_stamp", "timestamp")
                    .withColumnRenamed("temp_C", "temperature_celsius")
                    .withColumnRenamed("vibration_lvl", "vibration_level")
                    .withColumnRenamed("power_kW", "power_kw")
                    .withColumnRenamed("def_ct", "defect_count")
)
silver_sensor_df.columns

```

```
[11]: ['sensor_id',
       'machine',
       'timestamp',
       'temperature_celsius',
       'vibration_level',
       'power_kw',
       'defect_count',
       'status_flag',
       'noise_val',
       'extra_param']
```

Changing Types

```
[12]: silver_sensor_df.dtypes
```

```
[12]: [('sensor_id', 'string'),
       ('machine', 'string'),
       ('timestamp', 'string'),
       ('temperature_celsius', 'string'),
       ('vibration_level', 'string'),
       ('power_kw', 'string'),
       ('defect_count', 'string'),
       ('status_flag', 'string'),
       ('noise_val', 'string'),
       ('extra_param', 'string')]
```

```
[13]: silver_sensor_df = silver_sensor_df \
       .withColumn("timestamp", col("timestamp").cast(TimestampType())) \
       .withColumn("temperature_celsius", col("temperature_celsius").
       ↪cast(DoubleType())) \
       .withColumn("vibration_level", col("vibration_level").cast(DoubleType())) \
       .withColumn("power_kw", col("power_kw").cast(DoubleType())) \
       .withColumn("defect_count", col("defect_count").cast(IntegerType())) \
       .withColumn("noise_val", col("noise_val").cast(DoubleType())) \
       .withColumn("extra_param", col("extra_param").cast(DoubleType()))

silver_sensor_df.dtypes
```

```
[13]: [('sensor_id', 'string'),
       ('machine', 'string'),
       ('timestamp', 'timestamp'),
       ('temperature_celsius', 'double'),
       ('vibration_level', 'double'),
       ('power_kw', 'double'),
       ('defect_count', 'int'),
       ('status_flag', 'string'),
       ('noise_val', 'double'),
```

```
('extra_param', 'double']]
```

Checking Missing Values

```
[14]: # Count missing (NaN/null) values in each column
missing_values = silver_sensor_df.select(
    [count(when(col(c).isNull(), c)).alias(c) for c in silver_sensor_df.columns]
)
missing_values.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+
|sensor_id|machine|timestamp|temperature_celsius|vibration_level|power_kw|defect_count|status_flag|noise_val|extra_param|
+-----+-----+-----+-----+-----+-----+-----+
|      108|      116|      336|              256|              292|      312|
336|          84|      252|          324|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
```

```
[15]: # Dropping NaN/null values - removing rows which contain them
silver_sensor_df = silver_sensor_df.dropna()
missing_values = silver_sensor_df.select(
    [count(when(col(c).isNull(), c)).alias(c) for c in silver_sensor_df.columns]
)
missing_values.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+
|sensor_id|machine|timestamp|temperature_celsius|vibration_level|power_kw|defect_count|status_flag|noise_val|extra_param|
+-----+-----+-----+-----+-----+-----+-----+
|          0|          0|          0|              0|              0|          0|
0|          0|          0|              0|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
```

```
[16]: unique_machines = silver_sensor_df.select("machine").distinct()
unique_machines.show()
```

```
+-----+
|      machine|
+-----+
```

```

|   Conveyor-400|
|LaserCutter-900|
|   Grinder-800|
|   unknown|
|   CNC-Mill-200|
|   RobotArm-600|
|   Lathe-300|
| DrillPress-100|
|   Press-700|
| 3DPrinter-1000|
|   Turbine-500|
|           N/A|
+-----+

```

```
[17]: silver_sensor_df = silver_sensor_df.filter((silver_sensor_df.machine != "N/A")
↪      & (silver_sensor_df.machine != "unknown"))
```

```
[18]: unique_machines = silver_sensor_df.select("machine").distinct()
unique_machines.show()
```

```

+-----+
|   machine|
+-----+
|   Conveyor-400|
|LaserCutter-900|
|   Grinder-800|
|   CNC-Mill-200|
|   RobotArm-600|
|   Lathe-300|
| DrillPress-100|
|   Press-700|
| 3DPrinter-1000|
|   Turbine-500|
+-----+

```

Checking Duplicates

```
[19]: print(f"Number of rows before removing duplicates: {silver_sensor_df.count()}")
silver_sensor_df = silver_sensor_df.dropDuplicates()
print(f"Number of rows after removing duplicates: {silver_sensor_df.count()}")
```

```

Number of rows before removing duplicates: 198140
Number of rows after removing duplicates: 49535

```

Checking Anomalies

```
[20]: silver_sensor_df.describe().show()
```

```

+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
+-----+
|summary|          sensor_id|      machine|temperature_celsius|
vibration_level|          power_kw|      defect_count|status_flag|
noise_val|          extra_param|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
+-----+
|  count|          49535|          49535|          49535|
49535|          49535|          49535|          49535|          49535|
49535|
|  mean|          NULL|          NULL|
53.32816715453714|1.8233606540829692|21.787283940647985|0.1859695165034824|
NULL| 0.5002125971535277| 150.0378318360755|
| stddev|          NULL|          NULL|
3.9770878025493053|0.5092000216898719| 1.815363019171402|0.6327945751348232|
NULL|0.28787625796276245|29.162658232450056|
|  min|0e033d6a-4e2b-41e...|3DPrinter-1000|          36.41|
-0.59|          14.13|          0|          N/A|          0.0|
100.0|
|  max|          unknown| Turbine-500|          69.32|
4.18|          30.08|          3|          unknown|          1.0|
200.0|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
+-----+

```

```

[21]: negative_vibration_count = silver_sensor_df.filter(col("vibration_level") < 0).
      ↪count()
      print(f"Rows with negative vibration_level: {negative_vibration_count}")

```

Rows with negative vibration_level: 18

```

[22]: silver_sensor_df = silver_sensor_df.filter(col("vibration_level") >= 0) #↵
      ↪vibration levels for machines should not be negative in a physical sense ->↵
      ↪could be an error in data reading

```

Saving

```

[23]: silver_sensor_path = "data_lake/silver/sensor_data_silver/"
      os.makedirs(silver_sensor_path, exist_ok=True)

      silver_sensor_df.write.format("delta") \
        .mode("overwrite") \
        .save(silver_sensor_path)

```

```
print("Cleaned sensor data successfully stored in the Silver Layer.")
```

Cleaned sensor data successfully stored in the Silver Layer.

1.12.2 Production Data

Renaming Columns

```
[24]: bronze_production_df.columns
```

```
[24]: ['log_id',  
       'product_type',  
       'units_produced',  
       'defective_units',  
       'time_stamp',  
       'MachineName',  
       'remarks',  
       'operator_id',  
       'batch_info']
```

```
[25]: silver_production_df = (  
        bronze_production_df.withColumnRenamed("time_stamp", "timestamp")  
                               .withColumnRenamed("MachineName", "machine")  
    )  
  
    silver_production_df.columns
```

```
[25]: ['log_id',  
       'product_type',  
       'units_produced',  
       'defective_units',  
       'timestamp',  
       'machine',  
       'remarks',  
       'operator_id',  
       'batch_info']
```

Changing Types

```
[26]: silver_production_df.dtypes
```

```
[26]: [('log_id', 'string'),  
       ('product_type', 'string'),  
       ('units_produced', 'string'),  
       ('defective_units', 'string'),  
       ('timestamp', 'string'),  
       ('machine', 'string'),  
       ('remarks', 'string'),  
       ('operator_id', 'string'),
```

```
('batch_info', 'string')]
```

```
[27]: silver_production_df = silver_production_df \
      .withColumn("timestamp", col("timestamp").cast(TimestampType())) \
      .withColumn("units_produced", col("units_produced").cast(IntegerType())) \
      .withColumn("defective_units", col("defective_units").cast(IntegerType()))
silver_production_df.dtypes
```

```
[27]: [('log_id', 'string'),
      ('product_type', 'string'),
      ('units_produced', 'int'),
      ('defective_units', 'int'),
      ('timestamp', 'timestamp'),
      ('machine', 'string'),
      ('remarks', 'string'),
      ('operator_id', 'string'),
      ('batch_info', 'string')]
```

Checking Missing Values

```
[28]: # Checking for NaN/null values
silver_production_df.select(
    [count(when(col(c).isNull(), c)).alias(c) for c in silver_production_df.
     ↪columns]
).show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|log_id|product_type|units_produced|defective_units|timestamp|machine|remarks|operator_id|batch_info|
+-----+-----+-----+-----+-----+-----+-----+-----+
|      8|          4|          20|          16|          8|          12|          16|
4|          8|
+-----+-----+-----+-----+-----+-----+-----+-----+
|log_id|product_type|units_produced|defective_units|timestamp|machine|remarks|operator_id|batch_info|
```

```
[29]: # Removing NaN/null values
silver_production_df = silver_production_df.dropna()
silver_production_df.select(
    [count(when(col(c).isNull(), c)).alias(c) for c in silver_production_df.
     ↪columns]
).show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|log_id|product_type|units_produced|defective_units|timestamp|machine|remarks|operator_id|batch_info|
```



```
erator_id|batch_info|
+-----+-----+-----+-----+-----+-----+-----+
|      0|      0|      0|      0|      0|      0|      0|
0|      0|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
```

```
[30]: unique_machines = silver_production_df.select("machine").distinct()
unique_machines.show()
```

```
+-----+
|      machine|
+-----+
|  Conveyor-400|
|LaserCutter-900|
|   Grinder-800|
|      unknown|
|   CNC-Mill-200|
|  RobotArm-600|
|    Lathe-300|
|DrillPress-100|
|    Press-700|
|3DPrinter-1000|
|   Turbine-500|
|          N/A|
+-----+
```

```
[31]: silver_production_df = silver_production_df.filter((silver_production_df.
↪machine != "N/A") & (silver_production_df.machine != "unknown"))
```

```
[32]: unique_machines = silver_production_df.select("machine").distinct()
unique_machines.show()
```

```
+-----+
|      machine|
+-----+
|  Conveyor-400|
|LaserCutter-900|
|   Grinder-800|
|   CNC-Mill-200|
|  RobotArm-600|
|    Lathe-300|
|DrillPress-100|
|    Press-700|
|3DPrinter-1000|
```

```
|    Turbine-500|
+-----+
```

Checking Duplicates

```
[33]: print(f"Number of rows before removing duplicates: {silver_production_df.
        ↪count()}")
silver_production_df = silver_production_df.dropDuplicates()
print(f"Number of rows after removing duplicates: {silver_production_df.
        ↪count()}")
```

Number of rows before removing duplicates: 11900

Number of rows after removing duplicates: 2975

Checking Anomalies

```
[34]: silver_production_df.describe().show()
```

```
+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+-----+
|summary|          log_id|  product_type|  units_produced|
defective_units|      machine|          remarks|operator_id|
batch_info|
+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+-----+
|  count|          2975|          2975|          2975|
2975|          2975|          2975|          2975|          2975|
|  mean|          NULL|
NULL|125.15327731092437|22.805042016806723|          NULL|          NULL|
NULL|          NULL|
| stddev|          NULL|          NULL| 35.66397813052963|
7.910582247185824|          NULL|          NULL|          NULL|
NULL|
|  min|00130a22-b979-4b9...|3D Printed Part|          42|
-1|3DPrinter-1000|Adjusted machine ...|          OP10|Batch-3DPrinter-1...|
|  max|ffee8087-c73c-43d...|    Turbine Hub|          237|
50|    Turbine-500|          unknown|          OP20|          unknown|
+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+-----+
```

```
[35]: negative_defective_unit_count = silver_production_df.
        ↪filter(col("defective_units") < 0).count()
print(f"Rows with negative defective units: {negative_defective_unit_count}")
```

Rows with negative defective units: 15

```
[36]: silver_production_df = silver_production_df.filter(col("defective_units") >= 0)
      ↪# defective units cannot be negative -> error
```

Saving

```
[37]: silver_production_path = "data_lake/silver/production_data_silver/"
      os.makedirs(silver_production_path, exist_ok=True)

      silver_production_df.write.format("delta") \
          .mode("overwrite") \
          .save(silver_production_path)

      print("Cleaned production data successfully saved to the Silver Layer.")
```

Cleaned production data successfully saved to the Silver Layer.

1.12.3 Maintenance Data

Renaming Columns

```
[38]: bronze_maintenance_df.columns
```

```
[38]: ['maintenance_id',
      'machine',
      'maintenance_date',
      'maintenance_type',
      'duration_minutes',
      'cost',
      'operator',
      'notes']
```

```
[39]: silver_maintenance_df = bronze_maintenance_df \
      .withColumnRenamed("maintenance_date", "timestamp") \
      .withColumnRenamed("operator", "operator_id")
      silver_maintenance_df.columns
```

```
[39]: ['maintenance_id',
      'machine',
      'timestamp',
      'maintenance_type',
      'duration_minutes',
      'cost',
      'operator_id',
      'notes']
```

Changing Types

```
[40]: silver_maintenance_df.dtypes
```

```
[40]: [('maintenance_id', 'string'),
      ('machine', 'string'),
      ('timestamp', 'timestamp'),
      ('maintenance_type', 'string'),
      ('duration_minutes', 'int'),
      ('cost', 'double'),
      ('operator_id', 'string'),
      ('notes', 'string')]
```

```
[41]: silver_maintenance_df = silver_maintenance_df \
      .withColumn("timestamp", col("timestamp").cast(TimestampType())) \
      .withColumn("duration_minutes", col("duration_minutes").
      ↪cast(IntegerType())) \
      .withColumn("cost", col("cost").cast(DoubleType()))
silver_maintenance_df.dtypes
```

```
[41]: [('maintenance_id', 'string'),
      ('machine', 'string'),
      ('timestamp', 'timestamp'),
      ('maintenance_type', 'string'),
      ('duration_minutes', 'int'),
      ('cost', 'double'),
      ('operator_id', 'string'),
      ('notes', 'string')]
```

Checking Missing Values

```
[42]: silver_maintenance_df.select(
      [count(when(col(c).isNull(), c)).alias(c) for c in silver_maintenance_df.
      ↪columns]
      ).show()
```

```
+-----+-----+-----+-----+-----+-----+-----+
----+-----+
|maintenance_id|machine|timestamp|maintenance_type|duration_minutes|cost|operator_id|notes|
+-----+-----+-----+-----+-----+-----+-----+
----+-----+
|          0|      0|          0|          0|          0|      0|      0|
0|      0|
+-----+-----+-----+-----+-----+-----+-----+
----+-----+
```

```
[43]: unique_machines = silver_maintenance_df.select("machine").distinct()
unique_machines.show()
```

```
+-----+
```

```

|      machine|
+-----+
|  Conveyor-400|
|LaserCutter-900|
|   Grinder-800|
|   CNC-Mill-200|
|   RobotArm-600|
|     Lathe-300|
| DrillPress-100|
|     Press-700|
|  3DPrinter-1000|
|    Turbine-500|
+-----+

```

Checking Duplicates

```

[44]: print(f"Number of rows before removing duplicates: {silver_maintenance_df.
      ↪count()}")
silver_maintenance_df = silver_maintenance_df.dropDuplicates()
print(f"Number of rows after removing duplicates: {silver_maintenance_df.
      ↪count()}")

```

Number of rows before removing duplicates: 984

Number of rows after removing duplicates: 246

Checking Anomalies

```

[45]: silver_maintenance_df.describe().show()

```

```

+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
|summary| maintenance_id| machine|maintenance_type| duration_minutes|
cost|operator_id| notes|
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
| count|           246|           246|           246|           246|           246|
246|           246|           246|
| mean|           NULL|           NULL|           NULL|104.3170731707317|
570.1231707317074|           NULL|           NULL|
| stddev|           NULL|           NULL|
NULL|56.06116442151969|403.81712937623195|           NULL|           NULL|
| min|002999f5-c036-476...|3DPrinter-1000|           Emergency|           31|
164.19|           MT100|Addressed minor l...|
| max|fdf1f306-6ace-4d8...| Turbine-500|           Unscheduled|           235|
1497.8|           N/A|Unexpected shutdo...|
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+

```

Saving

```
[46]: silver_maintenance_path = "data_lake/silver/maintenance_data_silver/"
os.makedirs(silver_maintenance_path, exist_ok=True)

silver_maintenance_df.write.format("delta") \
    .mode("overwrite") \
    .save(silver_maintenance_path)

print("Cleaned maintenance data successfully saved to the Silver Layer.")
```

Cleaned maintenance data successfully saved to the Silver Layer.

1.12.4 Operator Dimension Table

```
[47]: dim_operator_path = "dim_operator.csv"

dim_operator_df = spark.read.format("csv") \
    .option("header", "true") \
    .option("inferSchema", "true") \
    .load(dim_operator_path)
dim_operator_df.show(5)
dim_operator_df.printSchema()
```

```
+-----+-----+-----+
|operator_id| operator_name|operator_type|
+-----+-----+-----+
|      MT100|   Allison Hill|   Maintenance|
|      OP12|   Noah Rhodes|   Production|
|      MT101|Angie Henderson|   Maintenance|
|      MT102| Daniel Wagner|   Maintenance|
|      OP18|Cristian Santos|   Production|
+-----+-----+-----+
only showing top 5 rows
```

```
root
 |-- operator_id: string (nullable = true)
 |-- operator_name: string (nullable = true)
 |-- operator_type: string (nullable = true)
```

```
[48]: # Save to Silver Layer as a Delta table
silver_operator_path = "data_lake/silver/dim_operator_silver/"
os.makedirs(silver_operator_path, exist_ok=True)

dim_operator_df.write.format("delta") \
    .mode("overwrite") \
    .save(silver_operator_path)
```

```
print("Operator Dimension Table successfully saved to the Silver Layer.")
```

Operator Dimension Table successfully saved to the Silver Layer.

1.13 Gold Layer – Business Insights

```
[49]: # Load Silver Layer tables
sensor_df = spark.read.format("delta").load("data_lake/silver/
↳sensor_data_silver")
production_df = spark.read.format("delta").load("data_lake/silver/
↳production_data_silver")
maintenance_df = spark.read.format("delta").load("data_lake/silver/
↳maintenance_data_silver")
operator_df = spark.read.format("delta").load("data_lake/silver/
↳dim_operator_silver")
```

1.13.1 Daily Sensor Metrics

```
[50]: # Convert timestamp to date for daily aggregation
sensor_metrics = sensor_df.withColumn("date", to_date(col("timestamp"))) \
    .groupBy("machine", "date") \
    .agg(
        avg("temperature_celsius").alias("avg_temp_celsius"),
        avg("vibration_level").alias("avg_vibration_level"),
        avg("power_kW").alias("avg_power_kW"),
        sum("defect_count").alias("total_defective_sensor_readings"),
        count(when(col("status_flag") == "Stopped", True)).
↳alias("downtime_count")
    ) \
    .orderBy(col("downtime_count").desc())
sensor_metrics.show()
```

```
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
|      machine|      date| avg_temp_celsius|avg_vibration_level|
avg_power_kW|total_defective_sensor_readings|downtime_count|
+-----+-----+-----+-----+-----+
+-----+-----+-----+
|   CNC-Mill-200|2024-11-28| 59.226999999999999| 1.8848000000000003|
22.1024|                                     2|          50|
|   Grinder-800|2024-10-22|          51.4114|
1.8628000000000002|21.371000000000002|                                     18|
50|
|   Press-700|2024-10-08| 52.927000000000001|          1.7162|
21.3132|                                     15|          50|
|   Turbine-500|2024-11-27| 55.721200000000001|
```

```

1.9028|21.992800000000003|          7|          50|
|   Conveyor-400|2024-10-24|          54.218|          1.6632|
21.2046|          8|          50|
|   CNC-Mill-200|2024-11-16|56.274800000000006| 1.8396000000000001|
21.5108|          2|          50|
|   Grinder-800|2024-10-02|48.940999999999995|          1.7226|
21.0508|          5|          50|
| 3DPrinter-1000|2024-11-06|          49.4552|
1.7806|21.417199999999998|          14|          50|
|  RobotArm-600|2024-12-29|54.657599999999995| 1.8774000000000002|
21.9312|          7|          50|
|LaserCutter-900|2024-10-23|          50.4344|
1.8115999999999999|20.758999999999997|          10|
50|
|   CNC-Mill-200|2024-10-28|54.379799999999996|          1.7482|
21.4448|          6|          50|
|   Press-700|2024-11-04|          55.8904|
2.0220000000000002|22.529600000000002|          9|
50|
|   Grinder-800|2024-12-11|          53.3012| 1.8163999999999998|
21.9822|          3|          50|
|LaserCutter-900|2024-10-04| 49.400799999999999| 1.9265999999999999|
20.8426|          10|          50|
|  RobotArm-600|2024-12-04| 56.787999999999999| 1.8628000000000002|
22.4544|          20|          50|
|LaserCutter-900|2024-12-03|          50.5732| 1.7462000000000004|
21.8658|          16|          50|
|   Grinder-800|2024-10-08| 48.546000000000001|          1.6454|
21.1152|          6|          50|
|   CNC-Mill-200|2024-11-21|          58.8006|
1.9422|21.852800000000002|          16|          50|
|   Grinder-800|2024-11-06|          52.839|
1.9054|21.755399999999998|          13|          50|
|   Grinder-800|2024-11-21| 54.219400000000001| 1.9432000000000003|
22.174|          7|          50|
+-----+-----+-----+-----+-----+
-----+-----+-----+
only showing top 20 rows

```

1.13.2 Daily Production Metrics

```

[51]: # Convert timestamp to date for daily aggregation
production_metrics = production_df.withColumn("date",
↳to_date(col("timestamp"))) \
    .groupBy("machine", "date") \
    .agg(

```



```

sum("units_produced").alias("total_units_produced"),
sum("defective_units").alias("total_defective_units"),
(sum("defective_units") / sum("units_produced")).alias("defect_rate"),
((sum("units_produced") - sum("defective_units")) /
sum("units_produced")).alias("production_yield")
) \
.orderBy(col("defect_rate").desc())
production_metrics.show()

```

```

+-----+-----+-----+-----+-----+
-----+-----+
|      machine|      date|total_units_produced|total_defective_units|
defect_rate| production_yield|
+-----+-----+-----+-----+-----+
-----+-----+
|  Turbine-500|2024-11-27|          114|
27|0.23684210526315788|0.7631578947368421|
|  Grinder-800|2024-11-13|          81|          19|
0.2345679012345679|0.7654320987654321|
|    Lathe-300|2024-10-15|          65|
15|0.23076923076923078|0.7692307692307693|
|    Press-700|2024-11-27|         265|
61|0.23018867924528302| 0.769811320754717|
|  Turbine-500|2024-12-02|         340|
78|0.22941176470588234|0.7705882352941177|
|  Turbine-500|2024-11-30|         315|
72|0.22857142857142856|0.7714285714285715|
|    Press-700|2024-12-19|         321|
73|0.22741433021806853|0.7725856697819314|
|DrillPress-100|2024-12-18|         285|
64|0.22456140350877193| 0.775438596491228|
|    Lathe-300|2024-11-17|         225|          50|
0.2222222222222222|0.7777777777777778|
|  Turbine-500|2024-12-04|         315|          70|
0.2222222222222222|0.7777777777777778|
|  Turbine-500|2025-01-03|         440|
97|0.22045454545454546|0.7795454545454545|
| Conveyor-400|2024-12-29|         992|
218|0.21975806451612903| 0.780241935483871|
|  Turbine-500|2024-12-07|         405|
89|0.21975308641975308|0.7802469135802469|
|    Lathe-300|2025-01-01|         456|
100|0.21929824561403508|0.7807017543859649|
|DrillPress-100|2024-12-19|         151|          33|
0.2185430463576159|0.7814569536423841|
|    Lathe-300|2024-10-29|         447|
97|0.21700223713646533|0.7829977628635347|

```

```

| RobotArm-600|2024-12-02|                272|
59|0.21691176470588236|0.7830882352941176|
|3DPrinter-1000|2024-12-02|                120|
26|0.21666666666666667|0.7833333333333333|
| Grinder-800|2024-12-11|                568|
123|0.21654929577464788|0.7834507042253521|
| Turbine-500|2024-12-28|                611|                132|
0.2160392798690671|0.7839607201309329|
+-----+-----+-----+-----+-----+
-----+-----+
only showing top 20 rows

```

1.13.3 Daily Maintenance Metrics

```

[52]: maintenance_metrics = maintenance_df.withColumn("date",
↳to_date(col("timestamp"))) \
    .groupBy("machine", "date") \
    .agg(
        sum("duration_minutes").alias("total_maintenance_duration"),
        count(when(col("maintenance_type") == "Scheduled", True)).
↳alias("scheduled_maintenance"),
        count(when(col("maintenance_type") == "Unscheduled", True)).
↳alias("unscheduled_maintenance"),
        count(when(col("maintenance_type") == "Emergency", True)).
↳alias("emergency_maintenance"),
        sum("cost").alias("total_maintenance_cost")
    ) \
    .orderBy(col("total_maintenance_cost").desc())
maintenance_metrics.show()

```

```

+-----+-----+-----+-----+-----+
-----+-----+
|      machine|      date|total_maintenance_duration|scheduled_maintenance|uns
cheduled_maintenance|emergency_maintenance|total_maintenance_cost|
+-----+-----+-----+-----+-----+
-----+-----+
|  Conveyor-400|2024-11-17|                209|                0|
0|                1|                1497.8|
|LaserCutter-900|2024-12-01|                216|                0|
0|                1|                1486.73|
|  RobotArm-600|2025-01-04|                176|                0|
0|                1|                1486.64|
|  RobotArm-600|2024-12-21|                179|                0|
0|                1|                1472.55|
|  RobotArm-600|2024-10-23|                214|                0|
0|                1|                1453.07|

```

3DPrinter-1000 2024-10-12	209	0
0 1	1443.28	
Turbine-500 2024-12-14	218	0
0 1	1414.64	
CNC-Mill-200 2024-12-27	167	0
0 1	1412.29	
RobotArm-600 2024-12-28	161	0
0 1	1411.01	
RobotArm-600 2024-12-12	138	0
0 1	1410.59	
Conveyor-400 2024-11-14	179	0
0 1	1390.82	
Conveyor-400 2024-11-05	150	0
0 1	1386.5	
Turbine-500 2024-11-19	180	0
0 1	1380.45	
LaserCutter-900 2024-11-18	212	0
0 1	1375.71	
Grinder-800 2024-10-27	174	0
0 1	1339.1	
CNC-Mill-200 2024-10-29	182	0
0 1	1329.23	
Conveyor-400 2024-10-02	140	0
0 1	1312.39	
Grinder-800 2024-11-10	148	0
0 1	1309.05	
Press-700 2024-11-11	192	0
0 1	1294.81	
DrillPress-100 2024-11-30	164	0
0 1	1291.8	

```

+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+

```

only showing top 20 rows

1.13.4 Joining Data

```

[53]: # Join sensor, production, and maintenance metrics on machine and date
gold_df = sensor_metrics \
    .join(production_metrics, ["machine", "date"], "left") \
    .join(maintenance_metrics, ["machine", "date"], "left")

gold_df.columns

```

```

[53]: ['machine',
      'date',
      'avg_temp_celsius',

```

```

'avg_vibration_level',
'avg_power_kW',
'total_defective_sensor_readings',
'downtime_count',
'total_units_produced',
'total_defective_units',
'defect_rate',
'production_yield',
'total_maintenance_duration',
'scheduled_maintenance',
'unscheduled_maintenance',
'emergency_maintenance',
'total_maintenance_cost']

```

1.13.5 Saving Data

```

[54]: save_path = "data_lake/gold/gold_machine_performance"
      os.makedirs(save_path, exist_ok=True)

      gold_df.write.format("delta").mode("overwrite").save(save_path)

```

```

[55]: gold_df = spark.read.format("delta").load(save_path)
      gold_df.show()

```

```

+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+
|      machine|      date| avg_temp_celsius|avg_vibration_level|      avg_pow
er_kW|total_defective_sensor_readings|downtime_count|total_units_produced|total_
defective_units|      defect_rate|  production_yield|total_maintenance_duratio
n|scheduled_maintenance|unscheduled_maintenance|emergency_maintenance|total_main
tenance_cost|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+
|   CNC-Mill-200|2024-11-07|57.463469387755104| 1.8671428571428568|
21.21081632653061|              3|              49|
175|              28|              0.16|              0.84|
NULL|              NULL|              NULL|              NULL|
NULL|
|   RobotArm-600|2024-11-21| 56.42632653061224|
1.8957142857142855|22.294693877551023|
13|              683|

```

128	0.18740849194729137 0.8125915080527086	NULL	
NULL		NULL	
	Press-700 2024-12-02 54.92416666666666 1.867499999999997	NULL	
22.456875		10	48 554
104	0.18772563176895307 0.8122743682310469	NULL	
NULL		NULL	
	Turbine-500 2024-11-27 55.72120000000001	NULL	
1.9028	21.992800000000003	7	50
114		27 0.23684210526315788 0.7631578947368421	
NULL		NULL	
NULL			
	Lathe-300 2024-10-26 52.358163265306125		
1.810816326530612	21.934489795918367	7	
49		110	
20	0.18181818181818182 0.8181818181818182	NULL	
NULL		NULL	
	3DPrinter-1000 2025-01-01 48.35918367346939		
1.8591836734693876	21.476938775510206	1	
15		139	
23	0.16546762589928057 0.8345323741007195	NULL	
NULL		NULL	
	3DPrinter-1000 2025-01-05 46.14459999999999 1.7688000000000001		
21.1772		1	21 267
42	0.15730337078651685 0.8426966292134831	NULL	
NULL		NULL	
	Conveyor-400 2024-10-24	54.218	1.6632
21.2046		8	50 395
58	0.1468354430379747 0.8531645569620253	NULL	
NULL		NULL	
	CNC-Mill-200 2024-11-16 56.274800000000006 1.8396000000000001		
21.5108		2	50 498
103	0.20682730923694778 0.7931726907630522	NULL	
NULL		NULL	
	LaserCutter-900 2024-11-15	52.0492	
1.9335999999999998	22.371399999999998	6	
7		130	
25	0.19230769230769232 0.8076923076923077	66	
1		0	0 339.83
	Grinder-800 2024-10-02 48.94099999999995	1.7226	
21.0508		5	50 446
70	0.15695067264573992 0.8430493273542601	NULL	
NULL		NULL	
	3DPrinter-1000 2024-11-06	49.4552	
1.7806	21.417199999999998	14	50
336		62 0.18452380952380953 0.8154761904761905	
NULL		NULL	
NULL			
	Lathe-300 2024-12-26 55.79104166666667		

2.0727083333333334	22.481458333333336		11
48	493		
98 0.19878296146044624	0.8012170385395537		NULL
NULL	NULL	NULL	NULL
LaserCutter-900 2024-12-17	51.4824		
1.7444 21.839799999999997		13	49
220	44	0.2	0.8
NULL	NULL	NULL	NULL
NULL			
3DPrinter-1000 2024-12-20	47.50163265306122		
1.7448979591836735 21.169183673469387			8
21	188		
25 0.13297872340425532	0.8670212765957447		NULL
NULL	NULL	NULL	NULL
RobotArm-600 2024-12-29	54.657599999999995	1.8774000000000002	
21.9312	7	50	707
145 0.2050919377652051	0.7949080622347949		NULL
NULL	NULL	NULL	NULL
LaserCutter-900 2024-10-23	50.4344		
1.8115999999999999 20.758999999999997			10
50	204		
30 0.14705882352941177	0.8529411764705882		NULL
NULL	NULL	NULL	NULL
Conveyor-400 2024-12-10	57.50755102040816	2.0106122448979593	
22.5234693877551		13	49
169	31 0.1834319526627219	0.8165680473372781	
NULL	NULL	NULL	NULL
NULL			
CNC-Mill-200 2024-10-28	54.379799999999996	1.7482	
21.4448	6	50	356
55 0.1544943820224719	0.8455056179775281		NULL
NULL	NULL	NULL	NULL
LaserCutter-900 2024-11-11	49.85166666666666		
1.76625 21.610833333333332		10	48
837	148 0.1768219832735962	0.8231780167264038	
NULL	NULL	NULL	NULL
NULL			
-----+	-----+	-----+	-----+
-----+	-----+	-----+	-----+
-----+	-----+	-----+	-----+
-----+	-----+	-----+	-----+
-----+	-----+	-----+	-----+


only showing top 20 rows

1.13.6 Visualizations

```
[56]: gold_pandas_df = gold_df.toPandas()
```

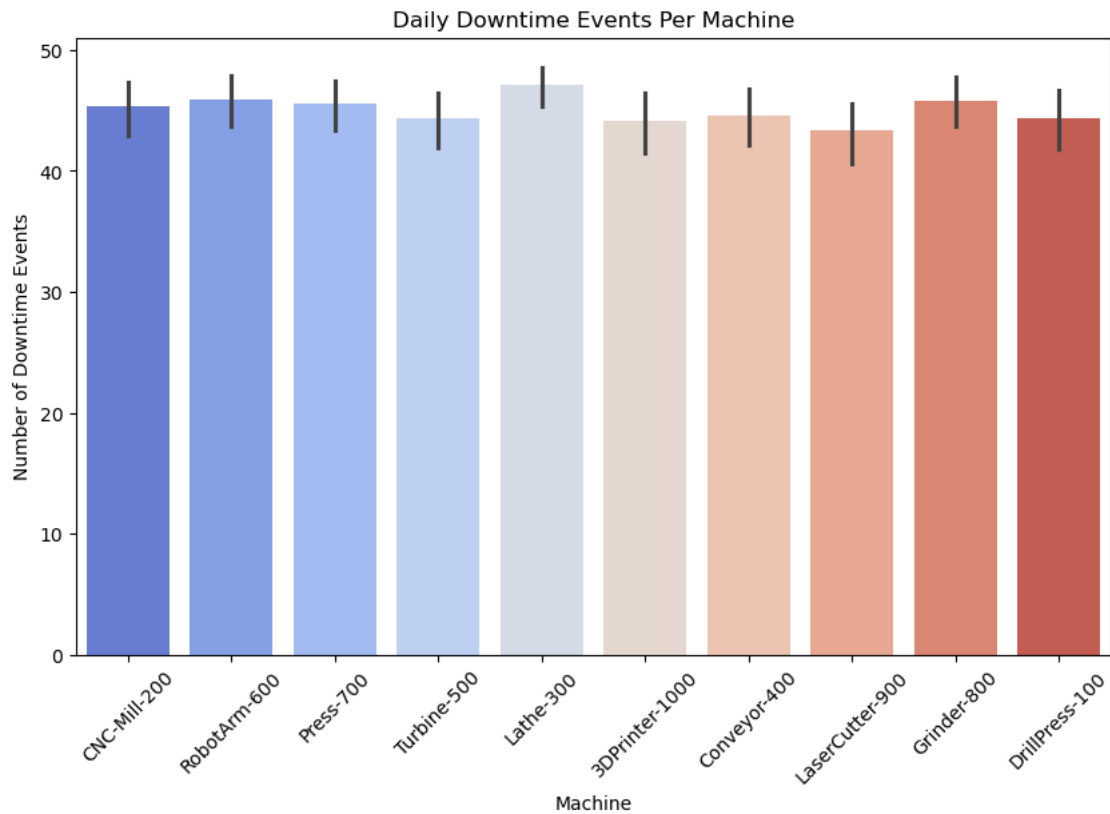
```
[57]: plt.figure(figsize=(12, 6))
sns.lineplot(data=gold_pandas_df, x="date", y="avg_temp_celsius",
             hue="machine", marker="o")
sns.lineplot(data=gold_pandas_df, x="date", y="avg_vibration_level",
             hue="machine", marker="s", linestyle="dashed")
sns.lineplot(data=gold_pandas_df, x="date", y="avg_power_kW", hue="machine",
             marker="^", linestyle="dotted")

plt.title("Daily Sensor Metrics (Temperature, Vibration, Power)")
plt.xlabel("Date")
plt.ylabel("Value")
plt.xticks(rotation=45)
plt.legend(title="Metric", labels=["Temperature (°C)", "Vibration Level",
                                   "Power (kW)"])
plt.show()
```



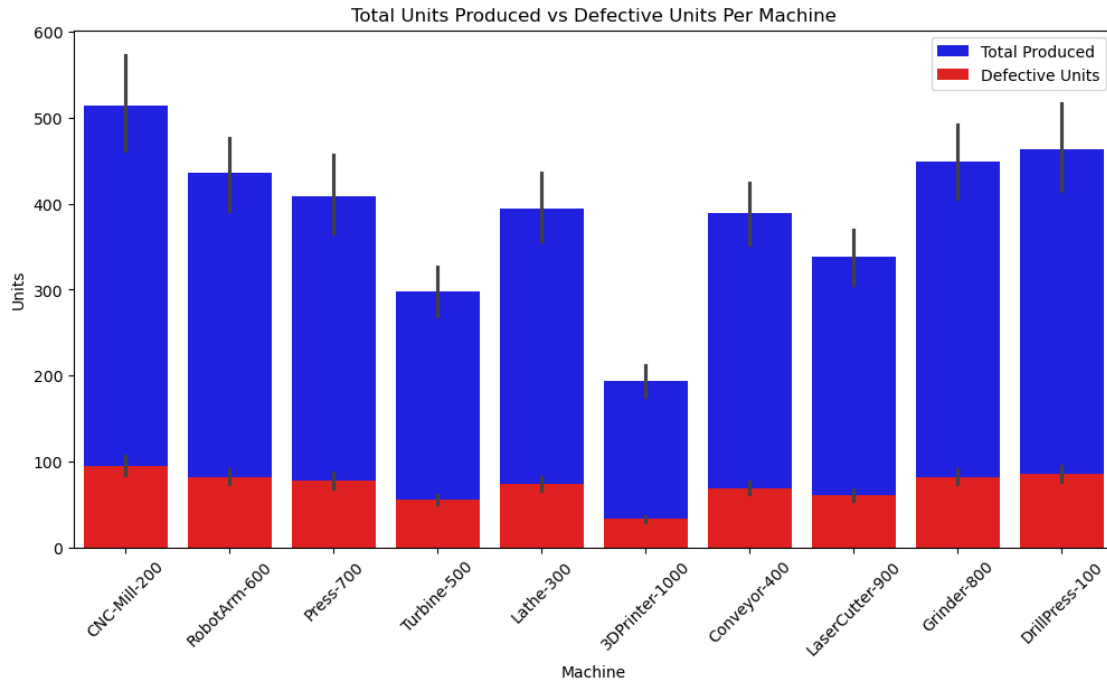
```
[58]: plt.figure(figsize=(10, 6))
sns.barplot(data=gold_pandas_df, x="machine", hue="machine",
            y="downtime_count", palette="coolwarm", legend=False)
plt.title("Daily Downtime Events Per Machine")
plt.xlabel("Machine")
```

```
plt.ylabel("Number of Downtime Events")
plt.xticks(rotation=45)
plt.show()
```

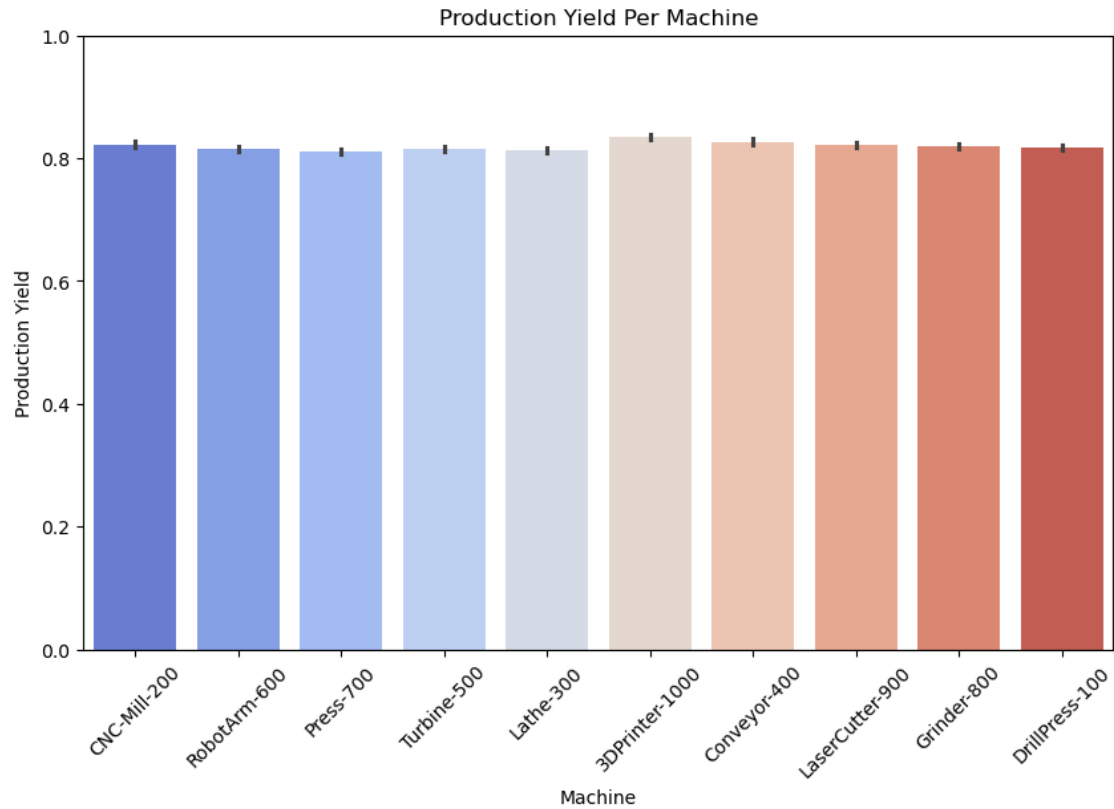


```
[59]: plt.figure(figsize=(12, 6))
sns.barplot(data=gold_pandas_df, x="machine", y="total_units_produced",
            color="blue", label="Total Produced")
sns.barplot(data=gold_pandas_df, x="machine", y="total_defective_units",
            color="red", label="Defective Units")

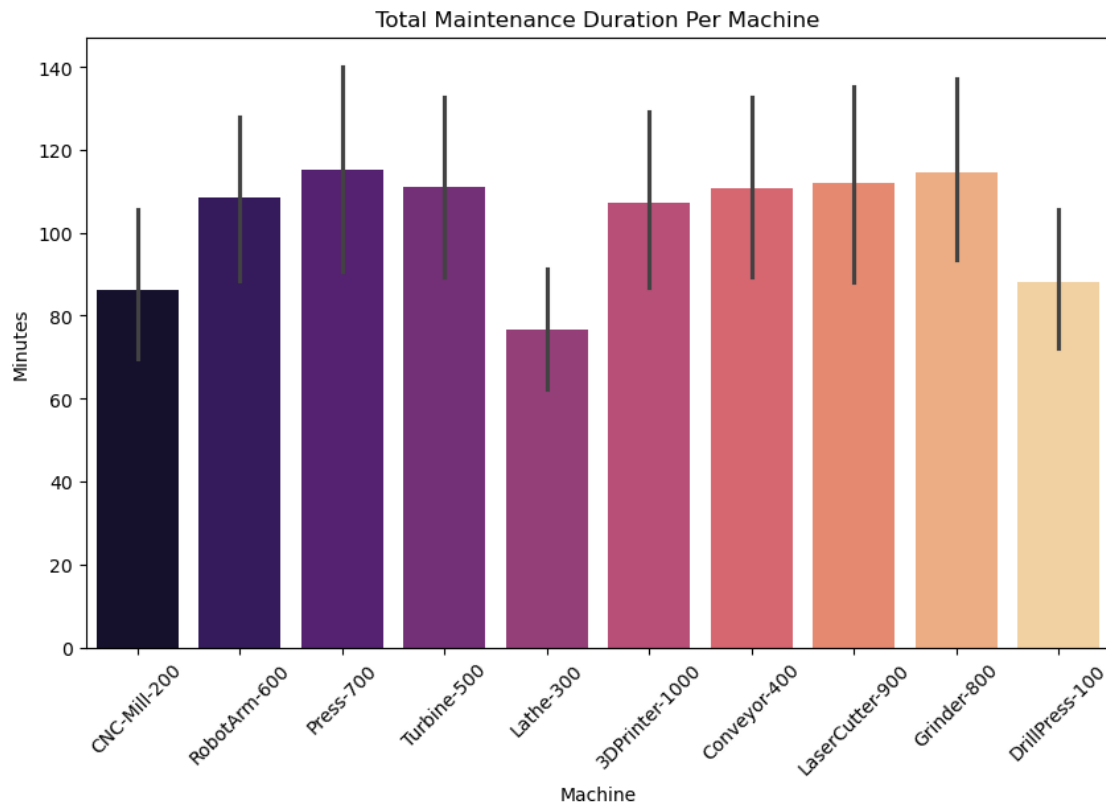
plt.title("Total Units Produced vs Defective Units Per Machine")
plt.xlabel("Machine")
plt.ylabel("Units")
plt.xticks(rotation=45)
plt.legend()
plt.show()
```

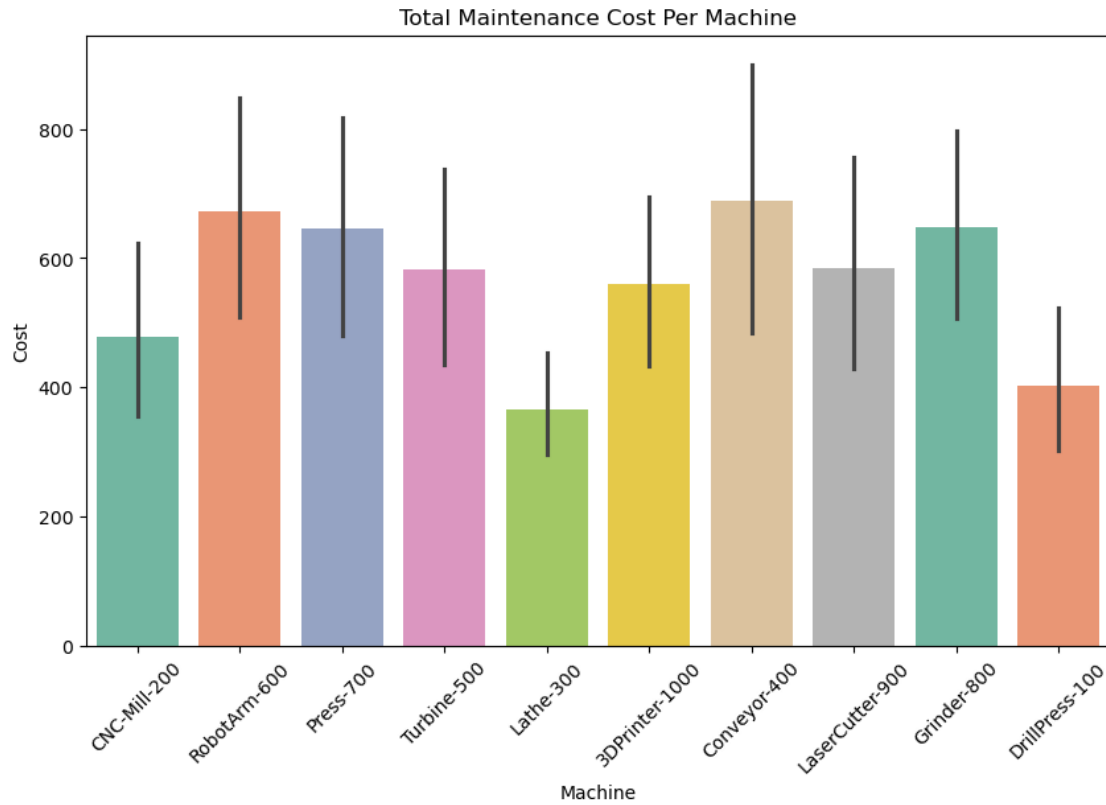
```
[60]: plt.figure(figsize=(10, 6))
sns.barplot(data=gold_pandas_df, x="machine", hue="machine",
            y="production_yield", palette="coolwarm")
plt.title("Production Yield Per Machine")
plt.xlabel("Machine")
plt.ylabel("Production Yield")
plt.ylim(0, 1)
plt.xticks(rotation=45)
plt.show()
```




```
[61]: plt.figure(figsize=(10, 6))
sns.barplot(data=gold_pandas_df, x="machine", hue="machine",
            y="total_maintenance_duration", palette="magma")
plt.title("Total Maintenance Duration Per Machine")
plt.xlabel("Machine")
plt.ylabel("Minutes")
plt.xticks(rotation=45)
plt.show()
```




```
[62]: plt.figure(figsize=(10, 6))
sns.barplot(data=gold_pandas_df, x="machine", hue="machine",
            y="total_maintenance_cost", palette="Set2")
plt.title("Total Maintenance Cost Per Machine")
plt.xlabel("Machine")
plt.ylabel("Cost")
plt.xticks(rotation=45)
plt.show()
```



```
[64]: plt.figure(figsize=(14, 6))
sns.scatterplot(data=gold_pandas_df, x="avg_vibration_level", y="defect_rate",
               hue="machine", size="total_units_produced", sizes=(20, 200), alpha=0.7)
plt.title("Defect Rate vs. Vibration Level")
plt.xlabel("Average Vibration Level")
plt.ylabel("Defect Rate")
plt.legend(title="Machine")
plt.show()
```



```
[65]: plt.figure(figsize=(12, 6))
sns.lineplot(data=gold_pandas_df, x="date", y="avg_power_kW", hue="machine",
             marker="o")
plt.title("Energy Consumption Trends Over Time")
plt.xlabel("Date")
plt.ylabel("Average Power (kW)")
plt.xticks(rotation=45)
plt.legend(title="Machine")
plt.show()
```



1.14 Discussion Questions (6p)

1. Have we now built a data lakehouse? Why or why not?

Yes, by looking at the characteristics of data lakehouses (unified-storage, ACID, scalability, integration with ML etc.) we can say that we have built a basic data lakehouse.

One aspect of lake houses is that they combine the flexibility of data lakes and reliability of data warehouses. We structured data into Bronze, Silver and Gold layers using Delta Lakes; thus, we ensured data consistency and schema enforcement. We then performed data cleaning and standardization in the Silver layer and aggregated business insights in the Gold layer. The use of Delta tables with ACID transactions ensured data integrity, which is one of the advantages Data lakehouses have over data lakes. We also performed automatic schema inference and schema evolution, which allows flexibility in handling changing data structures, which is one of the key aspects of delta lakehouses. Moreover, the dataframe we have in Gold layer is ready for Machine Learning pipelines.

Data lakehouses can handle both structured and semi- or unstructured data. Since we only had structured data, our implementation could not fully utilize the flexibility of a data lakehouse. Additionally, we did not implement advanced features like indexing or caching, which are common in lakehouse environments.

In overall, our implementation meets the core concepts of a data lakehouse, but it could be even extended to do advanced operations that data warehouses and data lakes are incapable of doing.

2. How does the medallion architecture enhance data quality and governance in this pipeline?

Bronze Layer: This layer stores raw data exactly as extracted from different sources. By keeping a copy of raw data, it preserves data lineage and enables us to recover from errors or reprocess data if needed. Since this layer acts as a historical archive, it ensures data traceability and compliance by maintaining the record of all incoming data. It also provides schema enforcement, preventing invalid records from corrupting the data lake.

Silver Layer: This layer is responsible for cleaning, standardizing, and validating the raw data. We performed the following transformations to improve data quality: - Renaming columns for consistency. - Removing duplicate records to eliminate redundancy. - Handling missing or corrupt values to ensure meaningful analysis. - Standardizing data types for compatibility across datasets. - Filtering out invalid sensor readings (e.g., negative values in vibration levels).

All these transformations make our data more structured and high quality. In terms of governance, this layer ensured schema evolution and validation, preventing incorrect data from propagating downstream.

Gold Layer: This layer aggregates and enriches the cleaned data to provide business insights. We created meaningful machine performance metrics such as: - Daily sensor metrics (e.g., average temperature, vibration levels, power usage). - Daily Production metrics (e.g., units produced vs. defective units, defect rates). - Daily maintenance metrics (e.g., analysis (total maintenance duration, cost per machine). - Advanced insight such as correlation between vibration levels and defects, energy consumption over time etc.

So this layer provides well-structured, curated datasets for analytics, reporting, and machine learning. It allows for fine-grained access control, ensuring that different user groups can access only relevant data while maintaining security and compliance.

3. What challenges might arise when scaling this pipeline from batch-based to real-time streaming data?

Scaling this pipeline from batch processing to real-time streaming might have some challenges. Data ingestion and latency become critical, requiring tools like Apache Kafka to handle continuous sensor, production, and maintenance data. Ensuring data quality is harder in a streaming setup, as records may arrive out of order or incomplete, requiring real-time validation and deduplication. Schema evolution also becomes a challenge since unexpected format changes can disrupt processing without proper handling. Additionally, fault tolerance and scalability must be addressed to handle system failures and spikes in data volume efficiently. Lastly, operational costs increase since streaming requires continuous resource usage.